Yaz vs prior knowledge-editing methods

Every method below is an editable model. Yaz's niche is the combination at <1M-parameter scale + explicit abstention—
not a new capability. ✓ = yes · ~ = partial / conditional · ✗ = no (author's fair reading of each method's core design).

	Yaz (2026)	ROME (2022)	MEMIT (2022)	GRACE (2023)	SERAC (2022)	PENME (2024)
Base model scale	<1M byte-level	1.5–6B	1.5–6B	0.1–1.5B	~0.4–1.5B	1–7B
Runs on CPU / offline	✓	✗	✗	~	~	✗
Update a fact (no retrain)	✓	✓	✓	✓	✓	✓
Delete a fact	✓	✗	✗	✓	✓	~
Add a new fact	✓	~	✓	✓	✓	✓
Paraphrase-robust routing	✓	✗	✗	~	✓	✓
Refuses when unsure (abstain)	✓	✗	✗	✗	✗	✗
Reversible edit / base frozen	✓	✗	✗	✓	✓	✓
Per-edit locality guarantee	✓	✗	✗	~	~	~
Open weights, pip-only repro	✓	~	~	~	~	~

Yaz: github.com/TilelliLab/Yaz · huggingface.co/TilelliLab/Yaz · MIT · routing 0.696 vs 0.216, abstention AURC 0.004 (oracle 0.003), 50 facts / CPU / seed 2026.

Honest scope: Yaz is a first-byte editor on a tiny synthetic benchmark; not state-of-the-art and not a moat. Refs: ROME/MEMIT (Meng+ 2022), GRACE (Hartvigsen+ 2023), SERAC (Mitchell+ 2022), PENME (2024, arXiv:2408.10411).